

How Does Google Train GBoard?

Introduction to Federated Learning

Google Keyboard

Next Word
Prediction

Out-of-
Vocabulary
Word Discovery

Emoji
Suggestions

Prerequisites

Basic Understanding of ML

Intermediate grasp over basic ML concepts and terminology

Desktop with Python IDE & A Python Environment

To extract the most out of this tutorial, it is recommended that you follow the coding exercises

Curiosity

Stay hungry, stay foolish - Steve Jobs

About Me

3rd Year
Computer
Science
Undergrad

Interned at
Acronis
Researching
Federated
Learning

Interested in All
Things AI/ML!

01

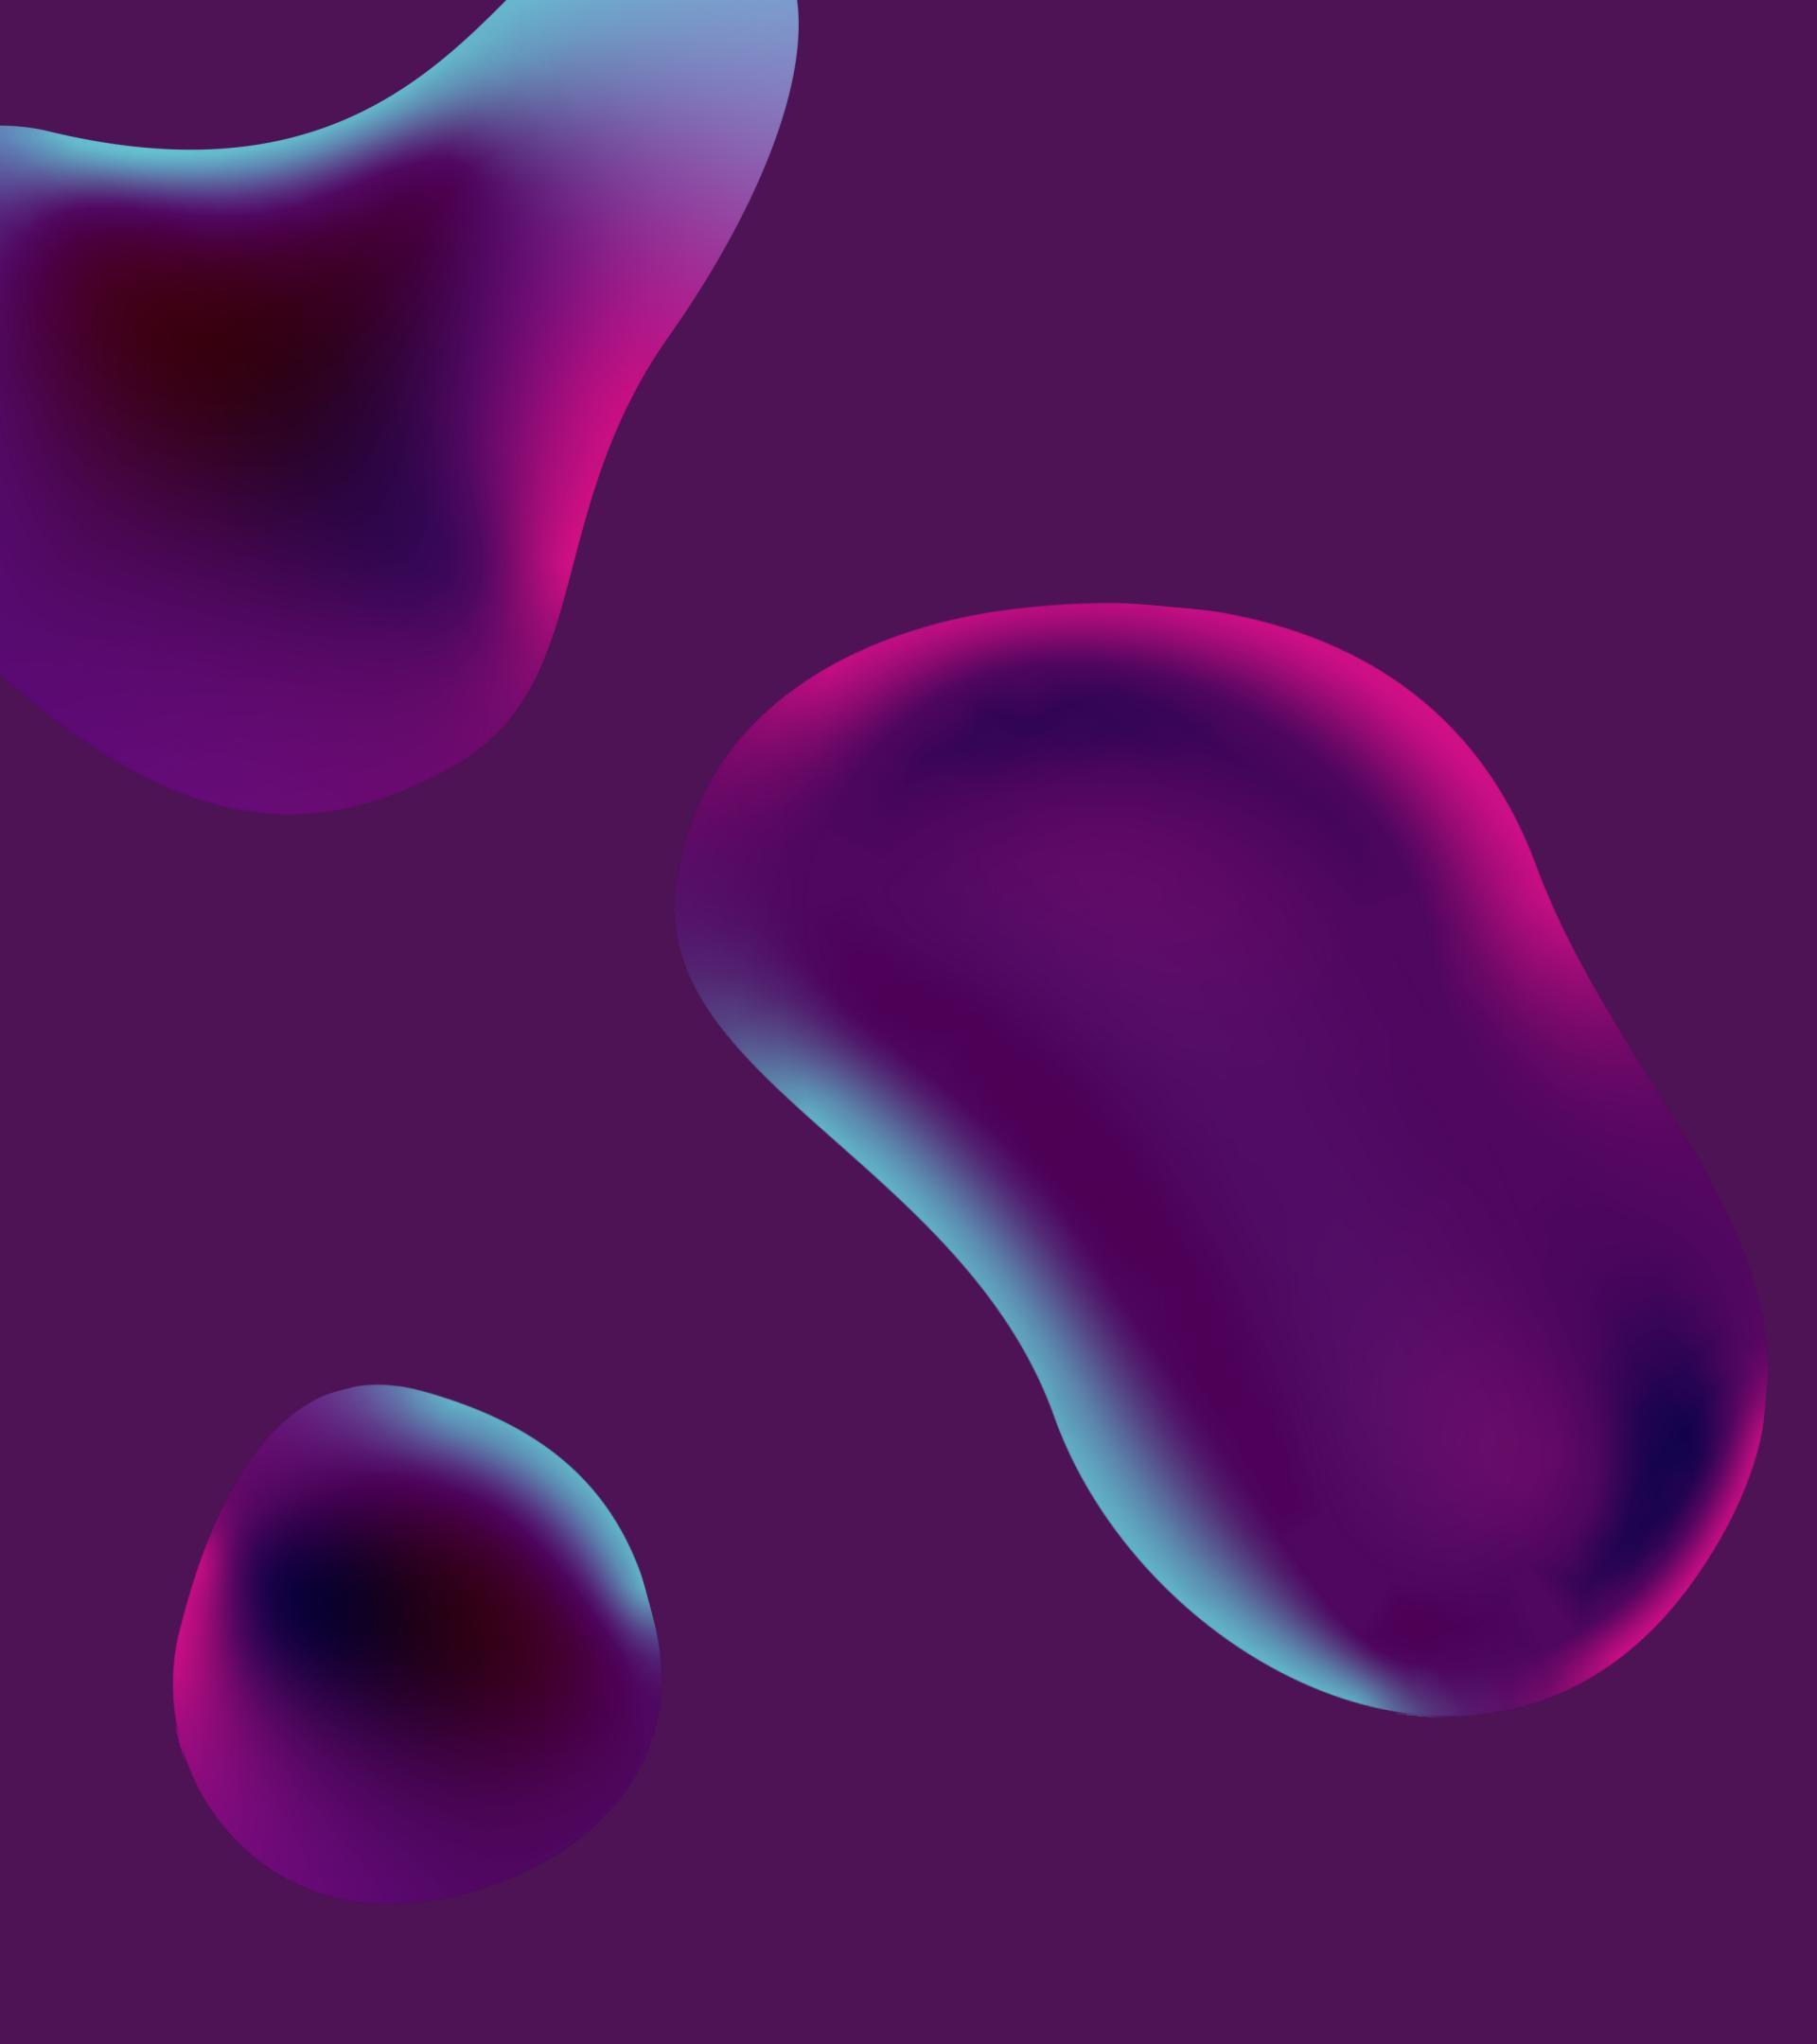
What is Traditional Machine Learning?

- All the data is stored in one storage unit or a cluster of storage units in close proximity
- All the data is accessible to the parties in control of the storage cluster
- Each machine learning (ML) model is trained with a full view of the entire dataset

Drawbacks?

- Data doesn't remain private
- All data must be aggregated in one place
 - Healthcare
 - Fintech
 - Insurance
 - Personal texts & images

02



The Solution? Collaborative Learning

also known as Federated Learning

Federated learning is decentralised machine learning wherein an algorithm is trained across multiple edge devices holding local data samples, without exchanging them

This means that each edge device trains a model locally on its local samples

All these models are then aggregated again and again over several rounds of local training to then create a combined model

As a result, a robust ML model can be trained without any party ever having to reveal its data

Types of Federated Learning

Horizontal Federated Learning

Datasets held locally by clients have the same features but different samples

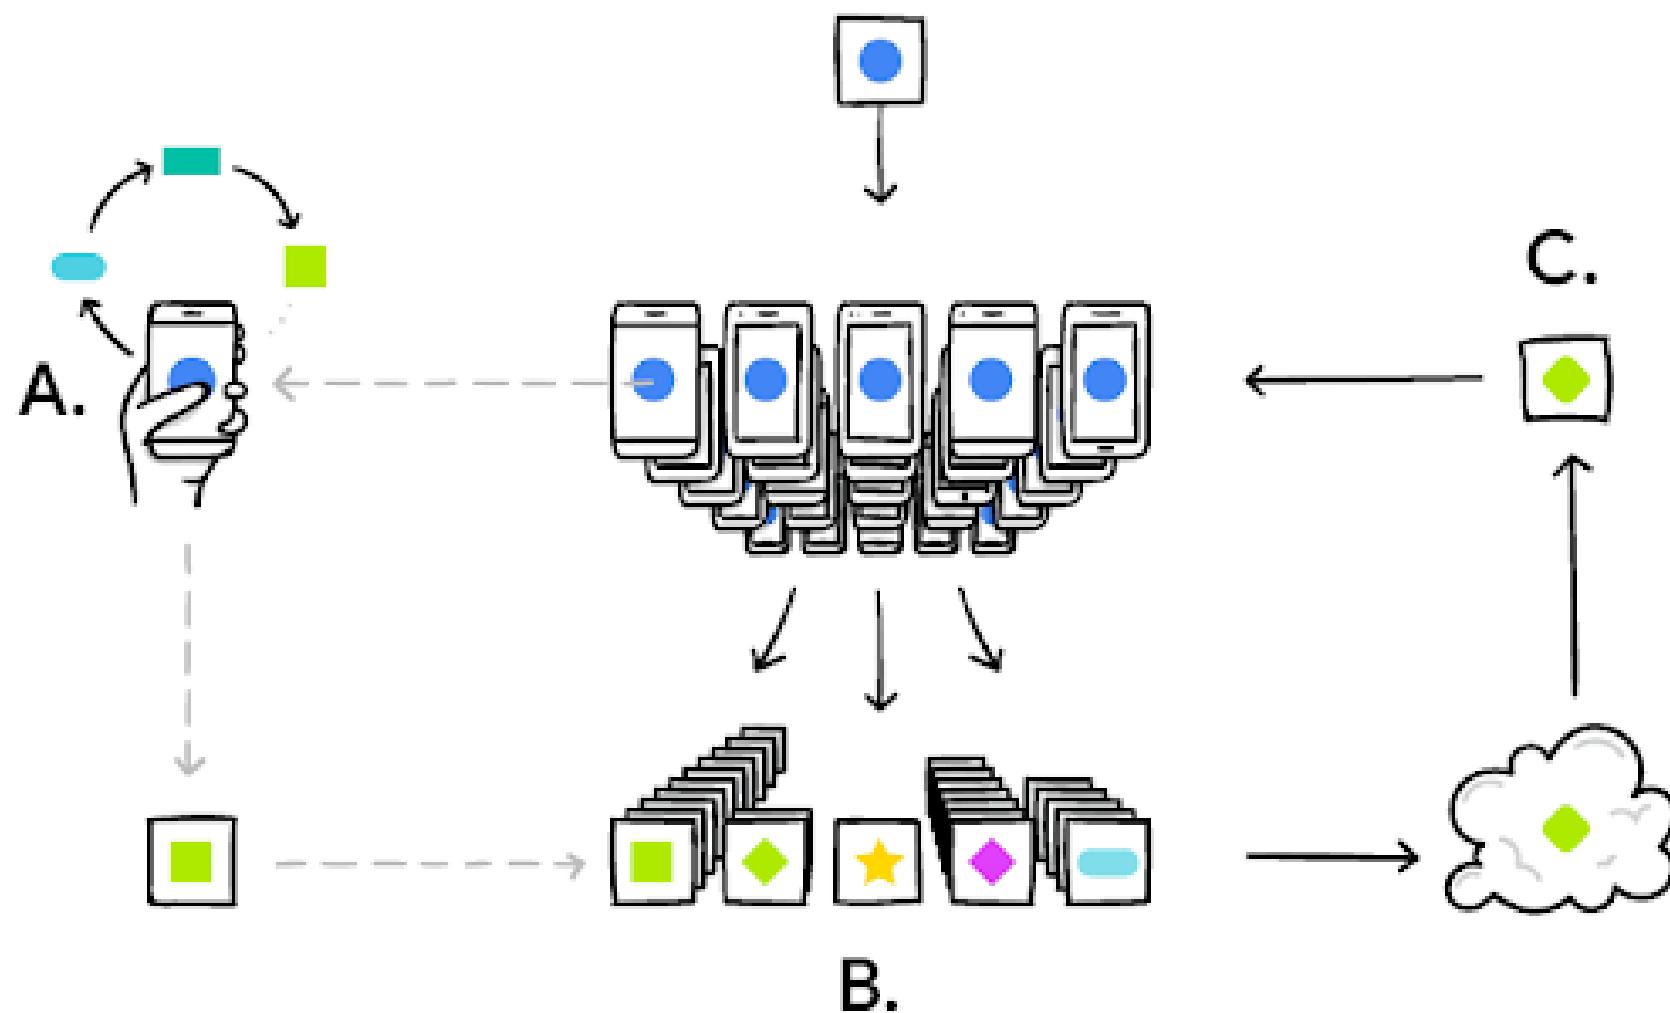
Vertical Federated Learning

Datasets held locally by clients have different features but the same samples

Federated Transfer Learning

Vertical FL utilised with a pre-trained model that is trained on a similar dataset for solving a different problem

How Exactly Does It Work?



Source: Google

A: Your phone builds a model personal to you

B: Everyone's phone also does the same for them

C: All these personalised models are aggregated to form a **consensus change** to the shared model

Back to A: Rinse & Repeat

What Happens On The Edge Device? (Client-side)

Receive Global Model

Each edge device receives the global model (randomly initialised) from the server

Update Model (Local Update)

The global model is updated by the edge device by performing Stochastic Gradient Descent (SGD) using its local samples

Return Updated Model

Finally, the edge device sends the updated model back to the server

Note: There might be several rounds of local updates (instead of just one) before the updated model is returned (Why?)

What Happens At The Server? (Server-side)

Initialise Global Model

The server randomly initialises a global model

Send Global Model

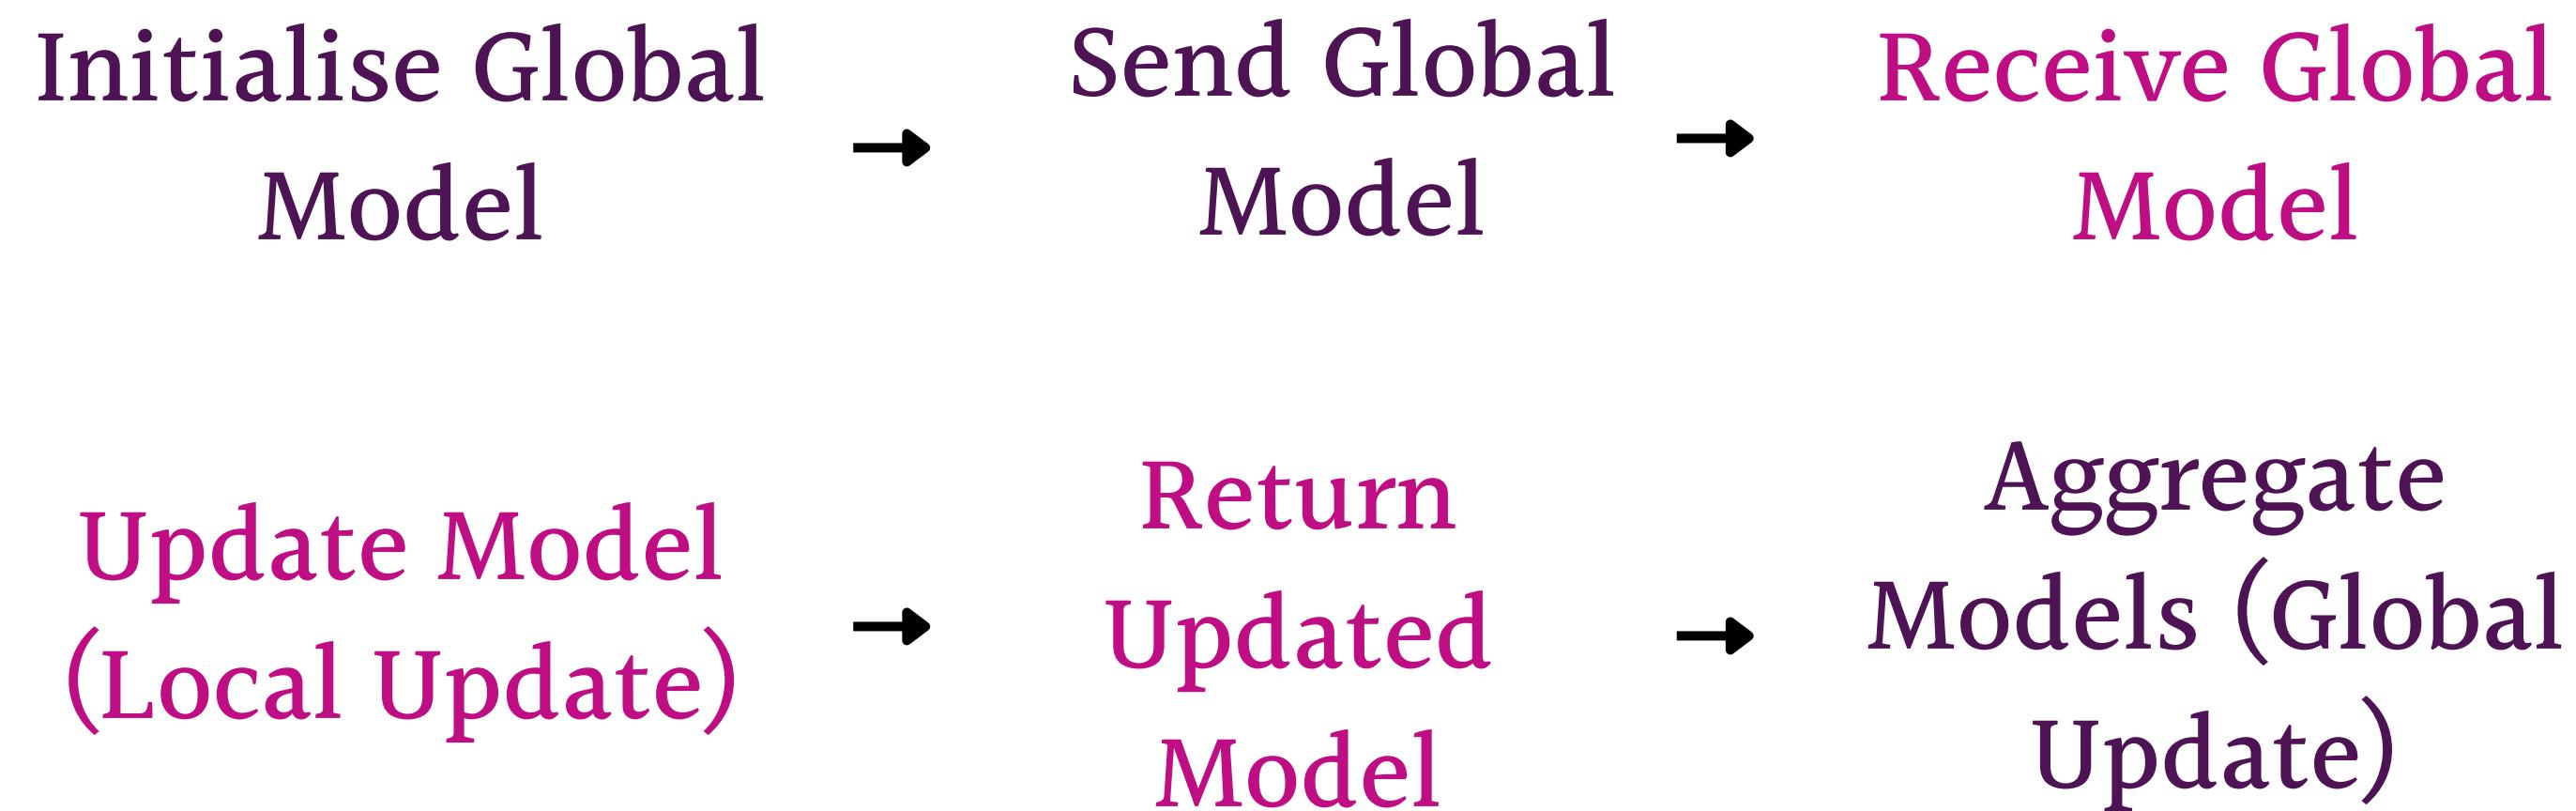
The server then picks a random subset of available clients to send the global model to

Aggregate Models (Global Update)

Finally, the server aggregates all the updated models sent back by the clients to update the global model

Note: The several different strategies used for aggregating the local updates (updated models) heavily influence the performance of the final model

Chronology



This cycle repeats for a predefined number of rounds

03

Federated Strategies

How **exactly** does the server combine the updated
models sent back by the clients?



intuitively, we can think of the local updates as one of two things

Updated Model

The client returns an **updated version of the model**

Therefore, the server receives a bunch of new (updated) models

Update to Model

The client returns how much the model should be updated by to reduce loss

Therefore, the server receives a bunch of values each indicating how **much** the model should change by



For this tutorial, we use the **2nd** definition. That is, each client calculates (and returns) the **average gradient** on its local data at the current model. Second definition is easier to implement. Why?

Terminology

w_t model weight at the beginning of round t

n_k number of local samples at client k

n total number of samples

η fixed learning rate (1 by default)

Federated Average

Client-side

$$g_k = \nabla F_k(w_t)$$

average gradient on its local data with the current model

Server-side

$$\nabla f(w_t) = \sum_{k=1}^K n_k/n * g_k$$

weighted average of all local
updates (local average gradients)

$$w_{t+1} = w_t - \eta \nabla f(w_t)$$

updated global model

Math Exercise

- In a federation of 3 clients, given an array of local gradient updates [13, 14.7, 15.92], an array of local sample sizes [250, 220, 175], and an initial model weight of 80.05, calculate the updated model weight after one round of federated learning. Assume all clients are chosen in every round and $\eta = 1$ (in practice, a very small subset of clients are chosen in any given round).
- What if you lower the value of η ?

Federated Average with Momentum

What is Momentum?

In FedAvg, the update weight of the model in **round $t+1$** depended solely on the gradients calculated in **round t**

However, with momentum, the updated weight of the model in **round $t+1$** depends not only on the gradients calculated in **round t** but also **round $t-1$, round $t-2$, ..., round 1**

The older the round, the lesser its impact; for example, by round 100, the impact of the gradients calculated in round 10 will be negligible, but still existent.

Federated Average with Momentum

Client-side

$$g_k = \nabla F_k(w_t) \longrightarrow \text{same as FedAvg}$$

Server-side

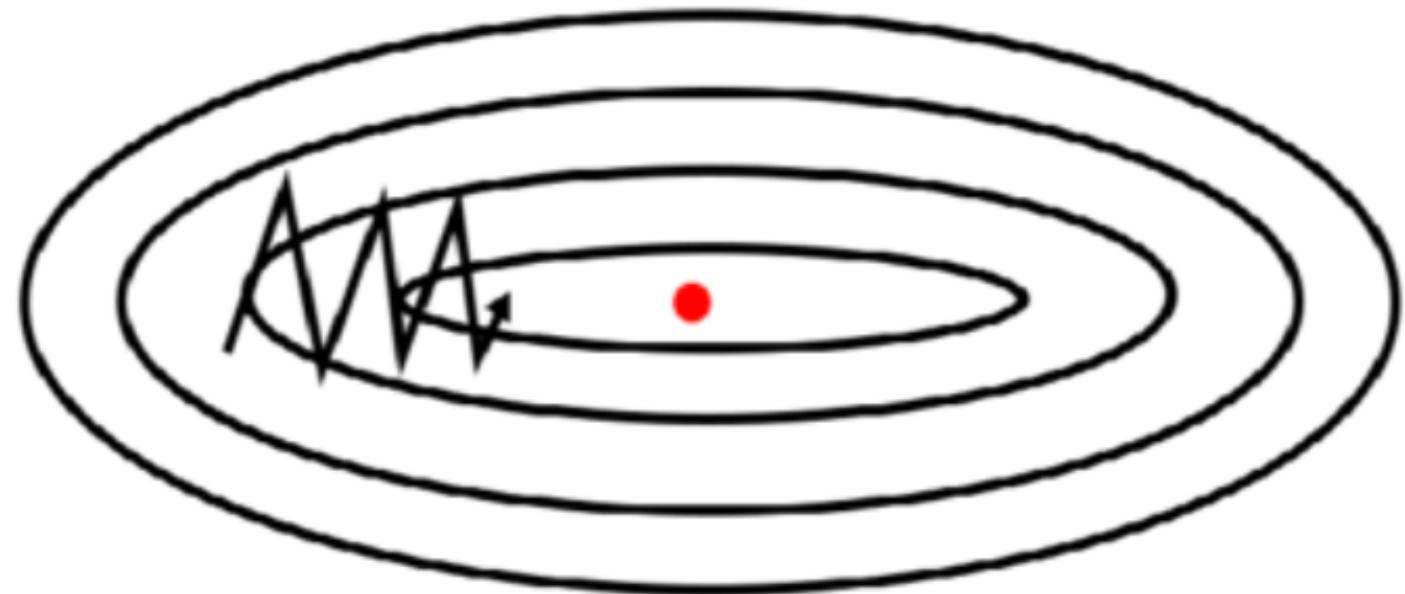
$$\nabla f(w_t) = \sum_{k=1}^K n_k/n * g_k \longrightarrow \text{same as FedAvg}$$

$$v_{t+1} = \beta v_t + \eta \sum_{k=1}^K (n_k/n) * g_k$$

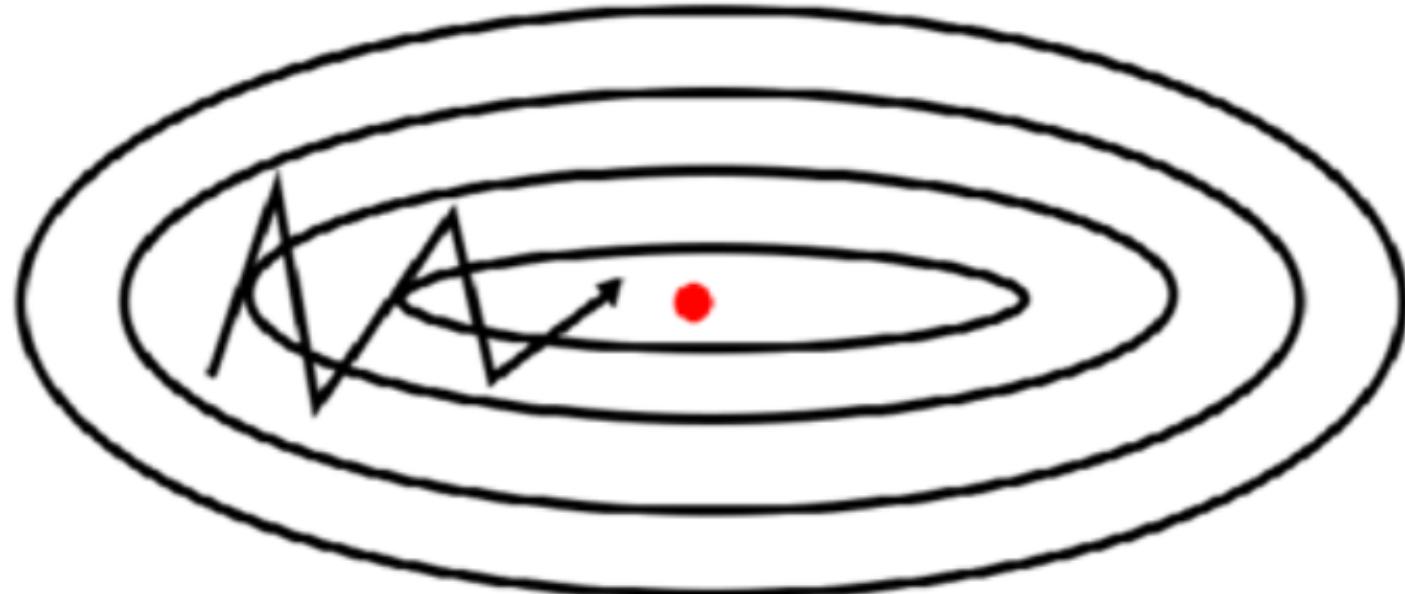
weighted average of all local weights discounted by a discounted factor of all previous weighted averages

$$w_{t+1} = w_t - v_{t+1}$$

SGD without momentum



SGD with momentum



Thinking Exercise

- Which strategy (out of FedAvg & FedAvgM) do you think fluctuates more? Why?
- What if you lower the value of beta?

Other Strategies

FedAdam

FedYogi

FedAdagrad

QFedAvg

FedDist

FedMa

FedPer

FaultTolerant

FedAvg

Caveats

Communication

The biggest bottleneck in FL is the communication between the client and the server

This is why there are often multiple rounds of SGD before local updates are sent back to the server—to minimise the back-and-forth and converge faster

Applicability

- Focused on non-convex neural network objectives
- However, is applicable to any finite-sum objective of the form $\min_{w \in R^d} f(w)$ where $f(w) = (1/n) * \sum_{i=1}^n f_i(w)$

Speedup

- Increase parallelism by adding more clients
- Has an upper limit

FL in Action

Let's now code a simple federated learning algorithm for a logistic regression model on a simple tabular dataset

We will compare the performance of traditional machine learning with both FedAvg and FedAvgM

04

Federated Learning Frameworks

What's Popular?

- Tensorflow Federated
- Flower
- Paddle Federated Learning
- PySyft
- OpenFL
- LEAF
- and many more

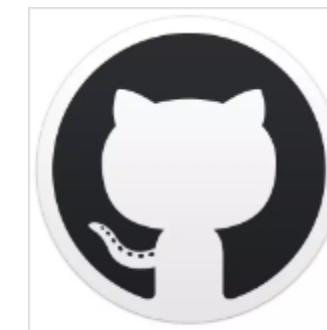
Why Flower?

- Compatible with most ML frameworks
- Allows deployment on edge devices such as phones
- Platform independent
- Easy to use

What Strategies Does it Support

- 7 most popular strategies
- Allows building and testing novel strategies

Code



architg1/Federated-Learning-Tutorial-DLW22-: The code and slides for my workshop on Federated Learning at MLDA...

The code and slides for my workshop on Federated Learning at MLDA Deep Learning Week 2022 - GitHub - architg1/Federated-Learning-Tutorial-DLW2...

 GitHub

<https://colab.research.google.com/gist/architg1/ecea3a3e21c0d6a97cae939e9c39da51/dlw-fl.ipynb>

<https://drive.google.com/drive/folders/13vLnzXGHDIZFngg8pLskA5u494fqHz1L?usp=sharing>

Installation

Python Env

```
conda install python=3.9.13
```

Flower

```
pip install flwr
```

Pandas

```
pip install pandas
```

Sklearn

```
pip install -U scikit-learn
```

Terminal (Mac)

```
pip uninstall grpcio  
conda install grpcio
```

wsl –install in an administrator command prompt (Windows)

Traditional ML

```
In [28]: from sklearn.linear_model import LogisticRegression
```

```
In [29]: model = LogisticRegression(  
        penalty="l2",  
        max_iter=1, # local epoch  
        warm_start=True, # prevent refreshing weights when fitting  
)
```

```
In [30]: model.fit(x_train, y_train)
```

Client-side

```
28     model = LogisticRegression(
29         penalty="l2",
30         max_iter=1, # local epoch
31         warm_start=True, # prevent refreshing weights when fitting
32     )
33
34     # Setting initial parameters
35     utils.set_initial_params(model)
36
37
38     class Client(fl.client.NumPyClient):
39         def get_parameters(self, config): # type: ignore
40             return utils.get_model_parameters(model)
41
42         def fit(self, parameters, config): # type: ignore
43             utils.set_model_params(model, parameters)
44             with warnings.catch_warnings():
45                 warnings.simplefilter("ignore")
46                 model.fit(x_train, y_train)
47             print(f"Training finished for round {config['server_round']}")
48             return utils.get_model_parameters(model), len(x_train), {}
49
50         def evaluate(self, parameters, config): # type: ignore
51             utils.set_model_params(model, parameters)
52             loss = log_loss(y_test, model.predict_proba(x_test))
53             accuracy = model.score(x_test, y_test)
54             return loss, len(x_test), {"accuracy": accuracy}
55
56
57     fl.client.start_numpy_client(server_address="0.0.0.0:8080", client=Client())
```

Server-side

```
22     # The `evaluate` function will be called after every round
23     def evaluate(server_round, parameters: fl.common.NDArrays, config):
24         # Update model with the latest parameters
25         utils.set_model_params(model, parameters)
26         loss = log_loss(y_test, model.predict_proba(x_test))
27         accuracy = model.score(x_test, y_test)
28         return loss, {"accuracy": accuracy}
29
30     return evaluate
31
32
33 # Start Flower server for three rounds of federated learning
34 if __name__ == "__main__":
35     model = LogisticRegression()
36     utils.set_initial_params(model)
37
38     # implementing 7 strategies
39     fed_avg_strategy = fl.server.strategy.FedAvg(
40         min_available_clients=2,
41         evaluate_fn=get_evaluate_fn(model),
42         on_fit_config_fn=fit_round,
43     )
44
45     fed_avg_m_strategy = fl.server.strategy.FedAvgM(
46         min_available_clients=2,
47         evaluate_fn=get_evaluate_fn(model),
48         on_fit_config_fn=fit_round,
49         server_learning_rate=1, # AssertionError: When using server-side optimization, model needs to be initialized.
50         server_momentum=0.00    # AssertionError: When using server-side optimization, model needs to be initialized.
51     )
52
53     fl.server.start_server(server_address="0.0.0.0:8080",
54                           strategy=fed_avg_strategy,
55                           config=fl.server.ServerConfig(num_rounds=3)
56                         )
```

What Can Go Wrong?

Hypothesis Conflicts

There is a hypothesis conflict between two classification models if there exists an input x such that the models output different classifications

Different Local Loss Surfaces

Each local model might reach to some local minima on their own surface, however, aggregated model might not be close to any local minima on the loss surface defined by the union of local datasets

How Do You Fix This?

Server-side Training

A small, balanced training dataset on server-side to fine-tune the aggregated results

Run SGD with Common Parameters

Since each local model is initialized to the same parameters, and each agent uses the same projection, we believe this might prevent local models from diverging too much from each other

05

Threats to Federated Learning

What Next?

Read the papers behind FedAvg,
FedAvgM, and FedOpt

Implement a federated strategy for
KNN or Random Forest using
Flower
Implement Federated Learning from
scratch in Python

References & Resources

<https://towardsdatascience.com/introduction-to-federated-learning-and-challenges-ea7e02f260ca>
<https://ai.googleblog.com/2017/04/federated-learning-collaborative.html>
<https://medium.com/bitgrit-data-science-publication/federated-learning-decentralized-ml-8709acfa9fa2>
<https://flower.dev/docs/quickstart-scikitlearn.html>
<https://flower.dev/docs/apiref-flwr.html#api-reference-flwr>
<https://arxiv.org/abs/2010.15582>
<https://arxiv.org/abs/2201.08135>
<https://towardsdatascience.com/ai-differential-privacy-and-federated-learning-523146d46b85>
<https://ai.googleblog.com/2022/02/federated-learning-with-formal.html>
<https://github.com/architg1/Federated-Learning-Tutorial-DLW22->